

Load Balancer with Service Turn-Off

API Design

Project statement

The project aims to enhance the functionality of a vanilla Kubernetes by enabling it to scale applications down to zero instances when they are not in use. This is relevant for ML applications that can be slow to start. The service will use event-driven automata to manage application scaling in response to real-time monitoring data and optimizing resource usage.

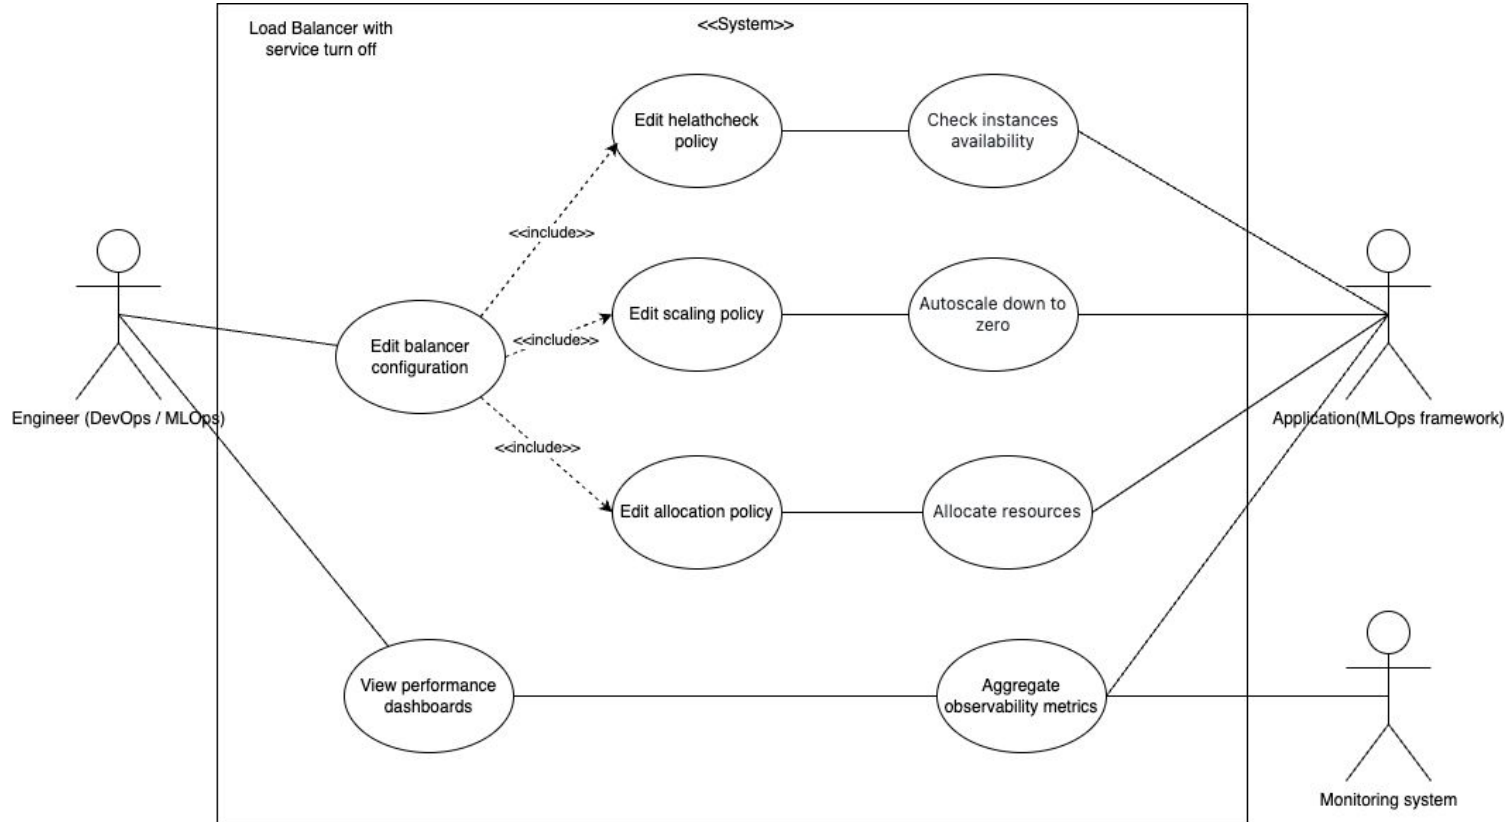
Team: Dmitry Kara, Daniil Mikulik, Ekaterina Karavayeva, Nikita Dumkin

Project repo: <https://github.com/dmitriykara/ads-tech-tornados-project>


This report:

<https://docs.google.com/presentation/d/1oxl-nbPDCdRBPtet2sjLLUGJs-bWyLksfpWnbOg8wT0/edit?usp=sharing>

Use case diagram



Service diagram



Open API

Using an Open API editor develop RESTful API for all microservices in the project, including data transfer objects. Store the API schema in your repository

(Note that the default example is excessive in detail)

<https://editor.swagger.io>

Verify step-by-step that your API supports scenarios or event flow

For each of the microservices show a scenario fragment/flow fragment where its API is invoked

<At least one microservice per team member>

API usage Policies Service

Policy ▾ {

id > [...]

type > [...]

spec > {...}

}

policies ^

POST /policies Create a new policy ▾

GET /policies/{id} Get a policy by ID ▾

PUT /policies/{id} Update a policy ▾

DELETE /policies/{id} Delete a policy by ID ▾

POST /policies/{id}/apply Apply a policy by ID ▾

API usage Node Manager Service

```
Node ▾ {  
  id* > [...]  
  title* > [...]  
  resources* > {...}  
  pods > [...]  
}
```

```
Pod ▾ {  
  id* > [...]  
  title > [...]  
  resources* > {...}  
}
```

node-manager

GET /nodes List all nodes

GET /nodes/{nodeId} Retrieve a node by ID

POST /nodes/{nodeId}/pods Create and allocate a new pod in a node

DELETE /nodes/{nodeId}/pods/{podId} Destroy a pod in a node

PUT /nodes/{nodeId}/pods/{podId} Reallocate a pod in a node

POST /nodes/{nodeId}/policy-update Push a signal about a policy update

API usage Monitoring Service

```
MonitoringEvent {  
  id* > [...]  
  payload* > [...]  
  observabilitySystemId* > [...]  
}
```

```
ObservabilitySystem {  
  id* > [...]  
  monitoringEventIds* > [...]  
  alertsIds* > [...]  
}
```

observability ^

POST**/events** Handle a new monitoring event**GET****/events/{eventId}** Retrieve a monitoring event by ID**POST****/observability-systems** Create a new observability system**GET****/observability-systems** List all observability systems**GET****/observability-systems/{observabilitySystemId}**
Retrieve an observability system by ID**PUT****/observability-systems/{observabilitySystemId}**
Update an observability system**DELETE****/observability-systems/{observabilitySystemId}**
Delete an observability system

API usage Balancer Service

BalancingAlgorithmRequest ▾ {

algorithmType > [...]
parameters > {...}

}

BalancingResult ▾ {

success > [...]
details > [...]

}

balancer ^

POST

/balancer/apply Применить алгоритм балансировки



POST

/balancer/events/process Обработать событие



POST

/balancer/events/send Послать событие об обновлении



Solution stack (prepare)

Find an example implementation of a microservices application in the programming language chosen.
Specify one value for each option below

Implementation

- OpenAPI
- Python, FastAPI
- RestAPI/JSON

Asynchronous interactions (optional)

- Celery

Testing tools

- Pytest

Operations

- Python3
- GitHub CI/CD
- Docker Compose
- Prometheus

Some references

<https://github.com/mfornos/awesome-microservices>

<https://awesomeopensource.com/projects/microservices-architecture>

<https://www.redhat.com/en/blog/comparing-openapi-grpc>

<https://cloud.google.com/apis/design/resources>